

Meeting Summary

Here's a summary of the group discussion:

The team is nearing the launch of a new product, with the final build for the launch candidate up on the staging server.

Key Activities & Decisions:

- **QA Testing:** Helly R conducted smoke tests and UI validations across multiple browsers, encountering a minor overflow issue on iPhone SE which was logged as low priority. All tests passed on Chrome and Edge, with Safari validation ongoing.
- **Pricing and Marketing Alignment:** Dylan G reminded the team to verify subscription tiering pages against the new pricing matrix. Mark S identified and fixed an issue with outdated pricing JSON. The final tagline, "Smart Chat, Local Mind," was approved. Marketing plans include teaser posts once QA signs off, with social posts scheduled for Monday morning (9 AM KST), and landing page draft approval pending.
- **Release Preparations:**
- Mark S planned to freeze dependencies at commit `a39f6` for the release branch.
- A code freeze was set for 11 PM that night, with no further merges after that.
- Helly R agreed to re-test the signup flow after the dependency freeze, reporting results before midnight.
- **User Feedback & Improvements:**
- Beta user feedback indicated satisfaction with speed but requested clearer onboarding.
- Dylan G proposed adding a tooltip walkthrough on first login, which Mark S implemented.
- Customer feedback highlighted some confusion around pricing tiers, leading Mark S to plan simplified pricing tooltips with example comparisons.
- **Technical Issues & Resolutions:**
- Helly R identified a delay on the notification socket when multiple users join simultaneously. Mark S planned to investigate and add debouncing.
- A critical issue with delayed push notifications on Android 14 was discovered. Helly R identified the cause as Firebase token expiration not being handled on reconnect and pushed a patch, which Dylan G redeployed.
- **Production Deployment & Monitoring:**
- The release candidate was successfully deployed to the production environment.
- The launch announcement was queued for Monday at 9 AM.

- The team will closely monitor metrics, with Helly R providing daily reports on signups and latency via a Grafana dashboard.
- Initial load testing showed stable server CPU under load.
- Analytics suggested pushing notifications during peak engagement hours (9-11 AM), leading to campaign messages being scheduled for those times with hashtags #LocalAI and #SmartChat.
- The team implemented autoscaling to handle a sudden surge in requests, with max instances set to 6.
- **Launch Success & Next Steps:**
- Uptime was reported as a steady 99.96%.
- Press coverage led to a significant traffic spike.
- The team plans to capture user testimonials from early adopters.
- Helly R will archive logs and prepare a metrics report for a Monday review.

Key People Involved:

- **Mark S:** Lead on build, deployment, technical fixes, and server stability.
- **Helly R:** Responsible for QA, testing, identifying technical issues, and monitoring.
- **Irving B:** Handles marketing, press, social media, and customer feedback.
- **Dylan G:** Focuses on product features, pricing alignment, analytics insights, and user outreach.